



DRMAA: Distributed Resource Management Application API

Andre Merzky

on behalf of


Peter Tröger

Hasso-Plattner-Institute (HPI) @ University of Potsdam

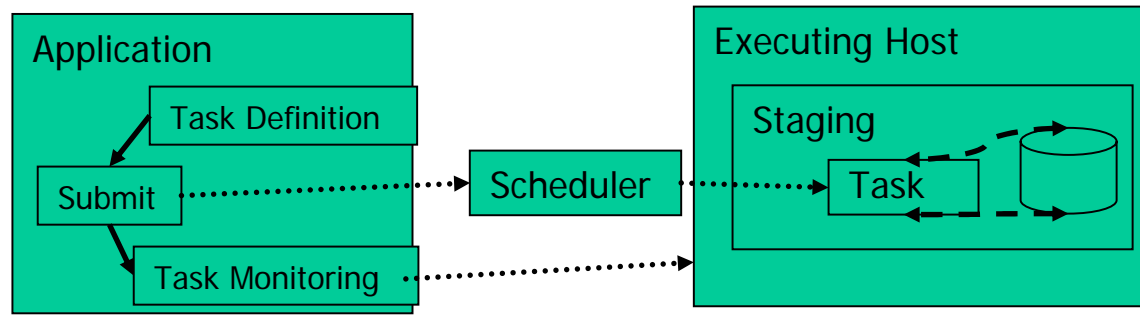
Scope: Run a Job API

(Steps from: "Ten Actions when SuperScheduling", GGF SchedWD 8.5, J.M. Schopf, July 2001)

- **Phase 1: Resource Discovery**
 - Step 1 Authorization Filtering
 - Step 2 Application requirement definition
 - Step 3 Minimal requirement filtering
- **Phase 2 System Selection**
 - Step 4 Gathering information (query)
 - Step 5 Select the system(s) to run on
- **Phase 3 Run job**
 - Step 6 (optional) Make an advance reservation
 - **Step 7 Submit job to resources**
 - Step 8 Preparation Tasks
 - **Step 9 Monitor progress (maybe go back to 4)**
 - **Step 10 Find out Job is done**
 - Step 11 Completion tasks



Resource Management Systems Differ Across Each Component



	Interface Format	Execution Environment
LSF	Has API plus Batch Utilities via “LSF Scripts”	User: Local disk exported System: Remote initialized (option)
Grid Engine	GDI API Interface plus Command line interface	System: Remote initialized, with SGE local variables exported
PBS	API (script option) Batch Utilities via “PBS Scripts”	System: Remote initialized, with PBS local variables exported
DataSynapse	Proprietary API.	User: Remote initialized

DRMAA Charter

- Develop an API specification for **the submission and control of jobs** to one or more Distributed Resource Management (DRM) systems.
- The scope of this specification is all the high level functionality which is necessary for an application to consign a job to a DRM system including **common operations on jobs** like termination or suspension.
- The objective is to **facilitate the direct interfacing of applications to today's DRM systems** by application's builders, portal builders, and Independent Software Vendors (ISVs).

DRMAA History

- BOF at GGF 3 in Frascati, Oct 2001
- WG status at GGF 4, Toronto, February 2002
- Participation from Altair (PBS), Sun Microsystems (SGE), Intel, IBM (LoadLeveler), University of Wisconsin (Condor), Cadence (Rocks system), Globus project
- Sideline engagement from EnFuzion, Entropia, Platform (LSF), GridIron project, United Devices
- June 2004: DRMAA 1.0 document accepted as *proposed recommendation* by GFSC
- Until today: Work on implementations and integration of user-provided feedback, new DRMAA compliance test suite

What have been the Issues?

- **General features**
 - Session concept
 - Asynchronous job monitoring
 - Scalability
 - Native features
- **Language bindings**
 - C/C++
 - Perl, Python
 - Fortran, Java
- **Libraries**
 - Serial / thread safe
 - Tracing / diagnosis
- **Advanced features**
 - Debugging support
 - File staging
 - Security
 - Job categories

Submit, control & monitor, and query status of jobs

DRMAA API Function Groups

- - Init / Exit
- - Job template handling
 - - Allocation / Deletion
 - - Job template parameter setter/getter routines
- - Job submission
 - - Individual jobs
 - - One time
 - - Multiple times – just re-adjust the job template (parameter sweep)
 - - Bulk jobs - implicit parameterization
- - Job monitoring and control
- - Auxiliary or system routines
 - - Error message routines
 - - Informational interfaces

Job Template

- - Description of all job requirements / parameters
- - Mandatory and optional parameters
- - Same intention as JSDL, but designed as ‘smallest common denominator’ between possible backend’s
- - Functions to create/delete job templates
 - `job_template *drmaa_allocate_job_template (void)`
 - `void drmaa_delete_job_template (job_template *jt)`
- - Setter/getter job template routines (scalar/vector)
 - `int drmaa_set_attribute (job_template *jt,`
 - `char *name,`
 - `char *value);`
 - `char* drmaa_get_attribute (job_template *jt,`
 - `char *name);`

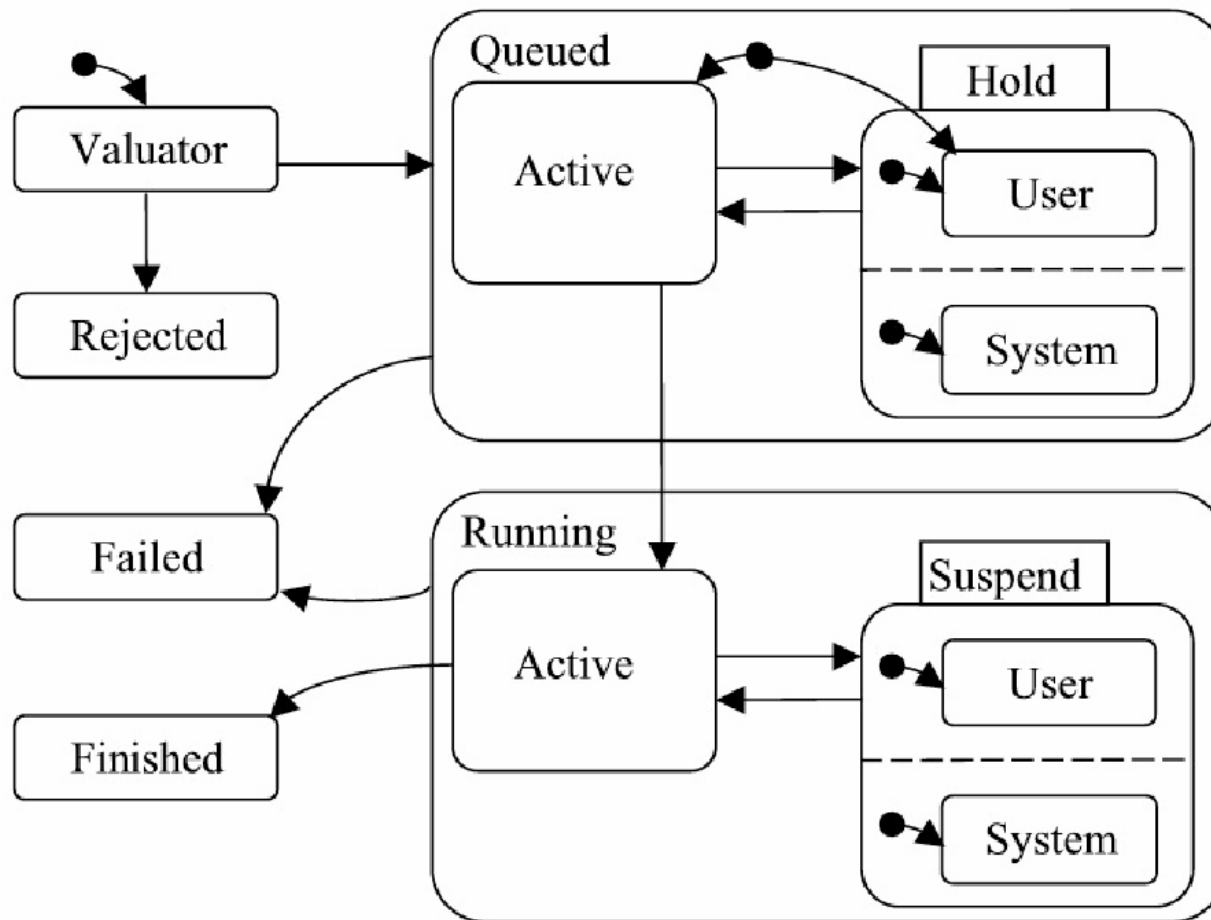
Job Submission

- - Jobs submitted to the DRM system are
- identified via an opaque job identifier (char*)
- - Single job identifiers are returned by
 - `int drmaa_run_job (job_template *jt,
char *job_id)`
- - Bulk job submissions return multiple job
- identifiers
 - `int drmaa_run_bulk_job (char **job_ids,`
 - `job_template *jt,`
 - `int start, int end, int incr);`

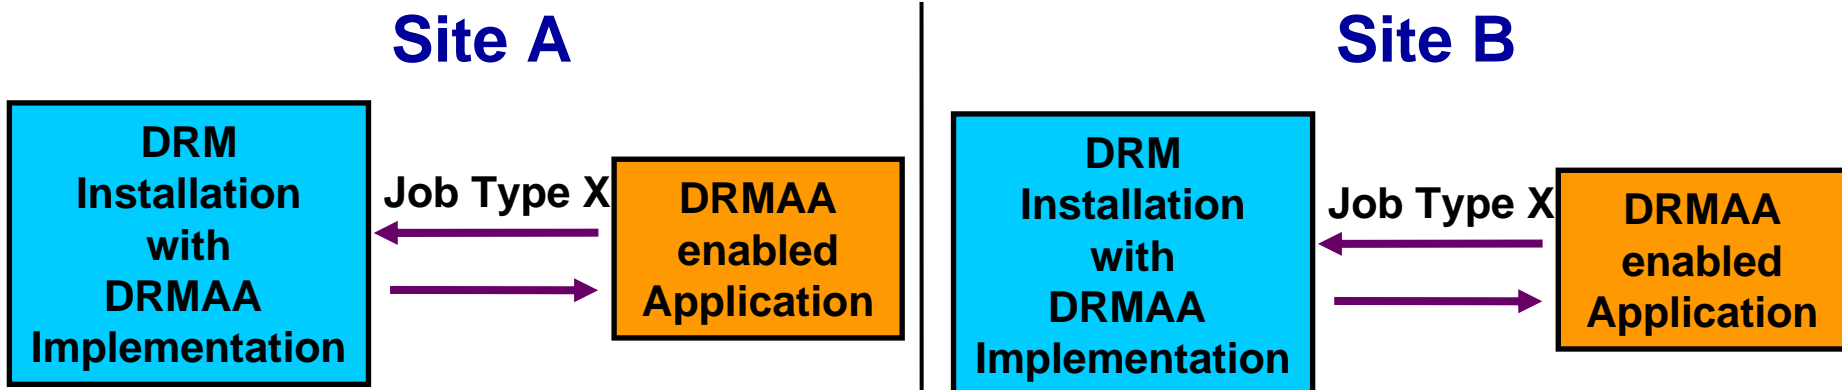
Job Monitoring, Control, and Status

- - Monitoring/Control functions
 - `int drmaa_control (char *job_id,`
`int action);`
 - `int drmaa_synchronize(char **job_ids);`
 - `int drmaa_job_ps (char *job_id,`
`int *remote_ps);`
- - Blocking and non-blocking waiting for one or
- more jobs to finish (like POSIX wait4(2))
 - `char *drmaa_wait(char *jobid, int *status,`
`int timeout, char **rusage);`
 - `drmaa_wif[exited|signaled|aborted]` and friends to get more information about failed jobs

Job State Transition



Job Categories



- Cluster consists of machines where X jobs run and others where they don't run

- X jobs run at all machines in cluster

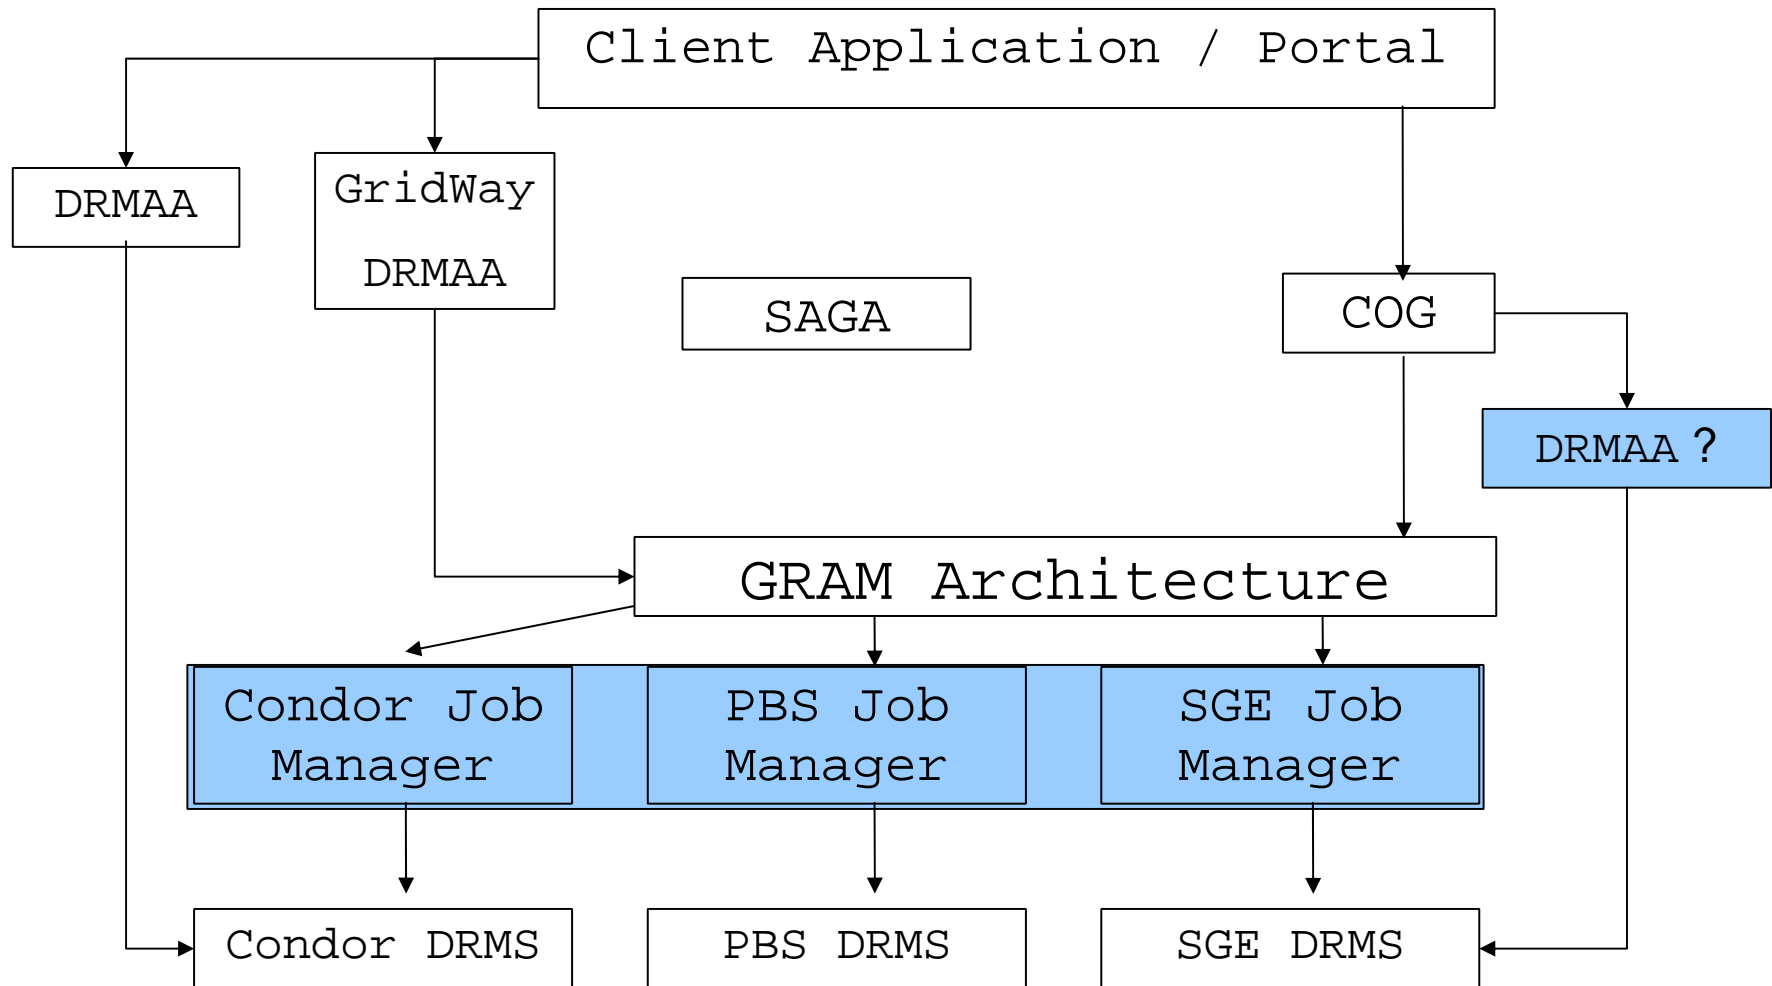
Native DRMS Options

- - The end user interacts with the DRMS via
- `native_specification` parameter
 - - Simple solution
 - - DRMAA implementation ignores the DRMAA DRMS
 - implicitly used and disallowed options
 - - Dist. Appls. Developers and DRMS vendors are not involved
 - in the local environment spec.
 - - The burden is on the end users to define the execution
 - environment
 - Need to know DRM
 - Need to know the remote application installation

DRMAA Placement

- - Realized:
 - - On top of DRM systems
 - - On top of Globus
 - - Beneath GRAM
 - - Interfaced by a Portal, application, shell
- - Also possible:
 - - UNICORE TSI interface to DRMSs
 - - CoG adapter
 - - On top of CoG
 - - Portable command line utilities (qsub, qstat)

A World of Submission API's



DRMAA in Practice

- - Multiple implementations since 2004
 - - Product implementation in Sun Grid Engine 6
 - - C- and Java-binding implementation
 - - Complete implementation in Condor 6.7 series
 - - C-binding implementation
 - - CPAN Perl DRMAA module (Tim Harsch)
 - - On-top-of DRMAA C-library
 - - Complete implementation in GridWay
 - - Allows DRMAA on-top-of Globus
 - - Prototype for Globus 3 DRMAA job manager (HPI)
 - - Based on DRMAA Perl implementation
- - Tutorials, programming examples, test suites
 - - <http://gridengine.sunsource.net>
 - - <http://www.dcl.hpi.uni-potsdam.de/research/drmaa>
 - - GGF12 tutorial, JavaOne 05 tutorial materials

Latest Improvements

- - Feedback from practical usage of available
- implementation(s) went back into spec
 - - Just look at the GridForge tracker and the SGE / Condor
 - mailing lists
- - Some details were under-specified
 - - Behavior in multi-threaded environments
- - C-centric style makes it hard to develop conformant
- object-oriented bindings
 - - Massive feedback from Java and .NET language binding
 - specification work
- - Several missing error codes fixed
- - But: Keep the API as small as it is !!!

The Future - DRMAA IDL Spec

- - Started work in early 2005
- - Based on Java- and .NET-binding experiences
- - Specification through standardized OMG Interface
- Definition Language (IDL)
 - - No, that does not mean CORBA ;-)
 - - Example: W3C DOM specification
 - - DRMAA language bindings will (and should) **not** rely on IDL
 - language bindings from OMG
 - - Complicated, weird semantics
 - - Simple custom binding by specifying consistent mapping rules
 - - Examples for Java in the IDL-spec
 - - Usage of IDL avoids wording issues (i.e. ,attribute‘ vs. ,property‘)
 - - Allows for true language-independent description of
 - namespaces, enumerations, constants, and time values

The Future - DRMAA IDL Spec (contd.)

- - Improved, more consistent description text for all
- functions
 - - More details regarding advanced OO-specific features
 - (multiple session objects, exception hierarchies)
 - - Consider languages with introspection functionalities
 - - Some details about RPC-DRMAA scenarios (SOAP, RMI, ...)
- - More parameter placeholders (e.g. for job ID)
- - A lot more possible error codes for the operations

The Future - Backward Compatibility

- - C- and Java bindings in their current state can be mostly
- derived also from the IDL spec
 - - Demand for consistent name mapping might change one or two
 - method names in the C-binding
 - - Introduction of new job state / error codes does not break
 - existing applications
 - - DRMAA has already a notion of versioning
- - .NET binding will be re-designed based on the IDL spec
- - No official binding documents for Perl and Python so far

DRMAA Documents and Tools

- - Updated DRMAA GFD-P-R
 - - Will be submitted to GFSG during the next weeks
 - - Three additional experience reports for Sun's N1GE, Condor and GridWay
- - DRMAA GFD-P-R or GFD.22 document
 - - Since June 2004
- - C binding v1.0
 - - Ready for submission to GFSG
- - Java binding 0.6.1
 - - Fairly complete
- - .NET binding v0.2
 - - Needs a sync
- - IDL document v0.35
 - - Nearly feature complete, needs to be augmented with latest DRMAA spec
 - - Will be submitted as a new, standalone GFD-P-R doc (might become DRMAA 2.0)
- - DRMAA Test Suite
 - - Hosted by Condor DRMAA CVS at Sourceforge: <http://sf.net/projects/condor-ext>
- - DRMAA Wiki
 - - Share your experience : <http://drmaa.org/wiki>

Next Steps

- - Submitting final document + experience reports
- - Continuous improvement of the test suite
- - Start work on next DRMAA version,
 - based on IDL spec
- - Collaboration with SAGA, JSDL, and OGSA-BES for
 - identifying synergy effects
- - More implementations, more languages, more
 - applications ...

Conclusion

- - Please take part in the discussion
 - Bi-weekly con calls
 - Toll Free: (866)545-5198 Code: 6898552
 - Regular: (865)521-8904
 - GridForge tracker
 - E-mail: *drmaa-wg@gridforum.org*
 - Archive: Use the link at *http://drmaa.org*
- - Please implement DRMAA and tell us your
- experience
 - It's easy, Dan did it 4 times ;-) ...

Thank you !